
PhysWM: Grounding Predictive World Models in Physical Equations

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Predictive world models can forecast dynamics without organizing their represen-
2 tations around the physical variables governing a system. We introduce PHYSWM,
3 a two-path architecture that grounds an action-conditioned predictive world model
4 through known physical equations. A neural path predicts the next state from data,
5 while an equation path maps the same predictive latent through a low-capacity
6 probe to physical variables and a frozen differentiable solver. Without parameter
7 labels, the equation path learns to reproduce the neural path’s stopped-gradient
8 prediction. We evaluate the method on FetchPush, randomized PushT, and Cart-
9 Pole. On FetchPush, the neural path beats persistence in every seed and equation
10 grounding reduces teacher-matching error from 8.758 to 0.752. On randomized
11 PushT, inferred effective dynamics beat shuffled episode assignments in all three
12 seeds. FetchPush exposes a solver-adequacy failure, while the small CartPole runs
13 remain below the data floor. These results support equation-compatible effective
14 coordinates without equating them with uniquely recovered material parameters.

15 1 Introduction

16 World models compress observations and actions into representations useful for predicting future
17 dynamics (4; 5; 6; 11). Predictive accuracy alone, however, does not ensure that these representations
18 expose mass, damping, stiffness, friction, or control response. A high-capacity predictor may exploit
19 trajectory correlations while leaving the governing physical factors entangled, even when its next-state
20 predictions are accurate (9). This separates predicting *what happens* from representing *why the system*
21 *evolves that way*.

22 We study whether known physical equations can ground a flexible predictive world model without
23 physical-parameter annotations. PHYSWM produces two next-state predictions from the same action-
24 conditioned latent (Figure 1). Path A is a learned decoder trained on observed transitions. Path B
25 maps the latent to a compact vector $\hat{\theta}$ and evaluates a frozen differentiable solver. Path B targets
26 Path A’s stopped-gradient prediction rather than the dataset state directly. The neural path remains
27 anchored to data and free to absorb model mismatch; the physical path asks whether the model’s
28 predictive belief admits a compact equation-based explanation. This fixed-decoder role is related
29 to differentiable system identification (8), but our training target is the predictor’s belief rather than
30 rendered pixels or labeled physical parameters.

31 Our contribution is threefold. First, we introduce a shared-latent, two-path objective for grounding
32 predictions in a supplied equation family without parameter labels. Second, we evaluate grounding
33 functionally by changing only the solver’s parameter vector while holding states, actions, weights,
34 and equations fixed. Third, we distinguish *semantic parameters*, which match named ground-truth
35 quantities, from *effective physical coordinates*, which improve equation rollouts but may compensate
36 for non-identifiability or solver mismatch.




Figure 1: **PHYSWM architecture.** Encoder and predictor form one shared trunk producing a single action-conditioned latent $\hat{z}$; both heads read $\hat{z}$, never a private view of the observation. Path A (blue) is the ordinary predictor: a decoder reads $\hat{z}$ to a state s_A , and $\mathcal{L}_A$ (Equation (4)) regresses the dataset’s ground truth, exactly what a predictive world model already does. Path B (orange) reads a physics vector θ off the same $\hat{z}$ with a deliberately low-capacity probe, integrates it through a *frozen, zero-parameter* solver (gray, dashed) to get an independent estimate s_B . $\mathcal{L}_B$ does *not* use the dataset label: it regresses s_A with gradient stopped (red, dashed), so the only route back to a learnable parameter is through the fixed solver into θ (orange dashed arrow), and θ can only improve by becoming a better physical explanation of what Path A already predicts. The dotted arrow is not part of the method: it marks the pre-action routing ablation (Section 3), where the probe reads z instead of $\hat{z}$.

2 Equation-grounded predictive representations

Given observations o_t , actions a_t , and prediction state s_t , an encoder and causal action-conditioned predictor form

$$z_t = E_{\phi}(o_t), \quad \hat{z}_t = P_{\psi}(z_{\leq t}, a_{\leq t}). \quad (1)$$

The neural decoder produces $\hat{s}_{t+1}^A = D_{\omega}(\hat{z}_t)$ and is trained in normalized state space:

$$\mathcal{L}_A = \|N(\hat{s}_{t+1}^A) - N(s_{t+1})\|_2^2. \quad (2)$$

A bounded low-capacity probe pools only causal context latents,

$$\hat{\theta} = \text{Bound}(\rho_{\xi}(\text{Pool}(\{\hat{z}_t\}_{t \in \mathcal{C}}))), \quad (3)$$

and a frozen solver S , containing no learnable parameters, produces $\hat{s}_{t+1}^B = S(s_t, a_t, \hat{\theta})$. The grounding loss is

$$\mathcal{L}_B = \|N(\hat{s}_{t+1}^B) - \text{sg}(N(\hat{s}_{t+1}^A))\|_2^2, \quad \mathcal{L} = \mathcal{L}_A + \alpha \mathcal{L}_B. \quad (4)$$

The stopped target prevents $\mathcal{L}_B$ from moving the Path-A decoder through its target edge, while gradients through S , ρ_{ξ} , and P_{ψ} shape the shared predictor. The default probe is linear, preventing it from hiding another world model. $\hat{\theta}$ is always a forward-pass output, never a free test-time optimization variable. The stop-gradient follows self-distillation methods (7; 1; 3); here it enforces a one-way dependency between a predictive teacher path and an equation-constrained student path.

Causal evaluation split. The identifier observes only completed context transitions. Metrics are computed on later query transitions, and no query outcome is accessible to the probe. With the default eight-frame window, four completed transitions form the context. Functional evaluation uses eight context transitions followed by seven held-out rollout transitions.

Grounding criterion. We call a coordinate physically grounded when reading it from the predictive latent allows the fixed equations to explain held-out dynamics better than a coordinate carrying no episode-specific information. This is deliberately weaker than claiming recovery of a unique material parameter: trajectories may identify only ratios such as stiffness/mass, and approximate solvers may induce compensating effective values.




Figure 2: Rendered observations from the three evaluated environments. FetchPush and CartPole use MuJoCo; PushT uses Pymunk contact dynamics.

Table 1: Mean $\pm$ sample standard deviation over three seeds. RMSE is lower better; R^2 is higher better.

Benchmark	Metric	Baseline / control	PHYSWM
FetchPush	Prediction RMSE	Persistence 0.648 ± 0.037	0.420 ± 0.015
FetchPush	Teacher-match RMSE	Ungrounded 8.758 ± 0.398	0.752 ± 0.036
PushT	Substitution RMSE	Shuffled 0.267 ± 0.013	Inferred 0.208 ± 0.011
CartPole	Substitution RMSE	Shuffled 0.339 ± 0.107	Inferred 0.311 ± 0.155
CartPole	Prediction RMSE	Persistence 0.199	0.792

3 Evaluation

We test four complementary properties. (1) *Predictive adequacy*: Path A must beat persistence, $\hat{s}_{t+1} = s_t$. (2) *Semantic readout*: where labels exist, a supervised linear probe measures held-out episode-level R^2 . (3) *Parameter substitution*: the solver receives the inferred vector, another episode’s shuffled vector, nominal constants, or true parameters when available. (4) *Multi-horizon rollout*: the same substitutions are rolled forward under fixed held-out actions. All errors are scaled per state dimension.

FetchPush. We use the Gymnasium-Robotics FetchPush task with per-episode object mass and friction variation. The MuJoCo environment supplies real contact dynamics, while the frozen solver uses a point-mass, linear-friction approximation. The reported primary runs use 256 episodes and three seeds.

Randomized PushT. This planar manipulation benchmark extends the contact-rich Push-T task (2) with episode-level variation in effective dynamics. Its solver is approximate and ground-truth labels are unavailable for most coordinates, making inferred-versus-shuffled substitution the primary test. The reported runs use 512 episodes and three seeds.

CartPole. The DeepMind Control Suite task (10) uses a six-parameter analytic solver for cart mass, pole mass, pole half-length, gravity, force gain, and damping. Only small, independently sampled runs of 8, 32, and 64 episodes are complete, so we treat CartPole as a pipeline and data-floor diagnostic rather than headline evidence.

4 Results and discussion

FetchPush satisfies the predictive premise. Path A reaches 0.420 ± 0.015 scaled RMSE, compared with 0.648 ± 0.037 for persistence, and wins in every seed (Table 1). The physics loss also changes the equation path substantially: its error against the stopped-gradient teacher falls from 8.758 ± 0.398 without grounding to 0.752 ± 0.036 with grounding. The shared latent therefore supports an equation-

PushT: ground truth vs PhysWM re-rendered predicted state, timestamp-by-timestamp pixel-diff error maps




Figure 3: Randomized PushT across consecutive timesteps. Rows show the ground-truth rendered observations, states predicted by PHYSWM and re-rendered through the environment, and their absolute pixel differences. Brighter pixels mark larger disagreement.




Figure 4: Randomized PushT substitution test over three seeds. Inferred episode-specific coordinates beat shuffled assignments in every seed and nominal parameters in two of three. Error bars show sample standard deviation.

82 constrained reconstruction of the neural prediction, even though the raw mass and friction coordinates
83 are not semantically identifiable at this data budget.

84 **PushT coordinates carry episode-specific dynamics.** Inferred variables obtain 0.208 ± 0.011
85 substitution RMSE versus 0.267 ± 0.013 for shuffled assignments (Fig. 4). They beat shuffled values
86 in all three seeds and nominal values in two of three. Solver sensitivity is concentrated in the agent’s
87 proportional and derivative gains, so this supports effective controller dynamics rather than complete
88 contact-parameter identification. Path A fails to beat persistence in one seed, which limits the strength
89 of this result.

90 **FetchPush coordinates compensate for solver mismatch.** The corrected frozen solver fails its
91 adequacy check even after directly fitting a free parameter vector to each episode: fitted solver RMSE
92 is 0.483, compared with 0.411 for persistence. In the functional run, inferred variables produce
93 lower error than true mass and friction (1.535 versus 6.984), while solver-space task success is
94 0.03125 for every parameter source. The probe is compensating for the solver’s fixed contact and
95 point-mass approximations. We therefore call $\hat{\theta}$ an effective coordinate rather than a recovered
96 material parameter.

97 **CartPole is still data-limited.** Across independent 8-, 32-, and 64-episode runs, inferred coordinates beat shuffled coordinates in a majority of seeds at every scale, but neither Path A nor the
 98 inferred solver beats its naive baseline. At 64 episodes, probe and shuffled substitution errors are
 99 0.311 ± 0.155 and 0.339 ± 0.107 , while nominal parameters reach 0.007. The non-monotonic Path-A
 100 error across budgets rules out a scaling claim from these small runs.
 101

102 5 Conclusion and limitations

103 PHYSWM grounds a predictive world model by requiring its shared action-conditioned representation
 104 to support a compact explanation through known physical equations. FetchPush establishes the
 105 predictive premise and shows that grounding makes the equation path track its teacher. Randomized
 106 PushT shows episode-specific information under parameter substitution. Neither result establishes
 107 recovery of uniquely named material parameters.

108 The study is limited to tiny-CNN diagnostics. The FetchPush solver is inadequate even after
 109 oracle fitting, PushT uses an approximate disc model for a T-shaped block, and CartPole has only
 110 small independent samples. The method does not discover equations, select informative actions,
 111 or demonstrate closed-loop planning or sim-to-real transfer. Larger CartPole runs, stronger visual
 112 encoders, and solvers that pass adequacy checks are required before broader claims.

113 References

- 114 [1] Mathilde Caron, Hugo Touvron, Ishan Misra, Hervé Jégou, Julien Mairal, Piotr Bojanowski,
 115 and Armand Joulin. Emerging properties in self-supervised vision transformers. In *Proceedings*
 116 *of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 2021.
- 117 [2] Cheng Chi, Siyuan Feng, Yilun Du, Zhenjia Xu, Eric Cousineau, Benjamin Burchfiel, and
 118 Shuran Song. Diffusion policy: Visuomotor policy learning via action diffusion. In *Proceedings*
 119 *of Robotics: Science and Systems (RSS)*, 2023.
- 120 [3] Jean-Bastien Grill, Florian Strub, Florent Altché, Corentin Tallec, Pierre H. Richemond, Elena
 121 Buchatskaya, Carl Doersch, Bernardo Avila Pires, Zhaohan Daniel Guo, Mohammad Gheshlaghi
 122 Azar, et al. Bootstrap your own latent: A new approach to self-supervised learning. In *Advances*
 123 *in Neural Information Processing Systems (NeurIPS)*, 2020.
- 124 [4] David Ha and Jürgen Schmidhuber. World models. *arXiv preprint arXiv:1803.10122*, 2018.
- 125 [5] Danijar Hafner, Timothy Lillicrap, Jimmy Ba, and Mohammad Norouzi. Dream to control:
 126 Learning behaviors by latent imagination. In *International Conference on Learning Representa-*
 127 *tions (ICLR)*, 2020.
- 128 [6] Danijar Hafner, Jurgis Pasukonis, Jimmy Ba, and Timothy Lillicrap. Mastering diverse domains
 129 through world models. *arXiv preprint arXiv:2301.04104*, 2023.
- 130 [7] Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. Distilling the knowledge in a neural network.
 131 *arXiv preprint arXiv:1503.02531*, 2015.
- 132 [8] Krishna Murthy Jatavallabhula, Miles Macklin, Florian Golemo, Vikram Voleti, Linda Petrini,
 133 Martin Weiss, Breandan Considine, Jerome Parent-Levesque, Kevin Xie, Kenny Erleben, Liam
 134 Paull, Florian Shkurti, Derek Nowrouzezahrai, and Sanja Fidler. gradsim: Differentiable
 135 simulation for system identification and visuomotor control. In *International Conference on*
 136 *Learning Representations (ICLR)*, 2021. Recovers mass, friction, and elasticity from video by
 137 backpropagating through a differentiable physics engine and a differentiable renderer end to
 138 end. Our solver plays an analogous fixed-decoder role, but is frozen rather than optimized, and
 139 the target it is trained against is a predictor’s own next-state belief rather than pixels.
- 140 [9] Zhenjiang Mao, Mrinal Eashaan Umasudhan, and Ivan Ruchkin. Physically interpretable world
 141 models via weakly supervised representation learning. *arXiv preprint arXiv:2412.12870*, 2024.
 142 Closest prior work in spirit: aligns latent states with physical quantities and constrains their
 143 evolution with a partially-known dynamics model, without ground-truth physical annotations.
 144 The supervision differs in kind: PIWM uses weak, distribution-based physical labels and a

- 145 learnable dynamics model; our training signal uses no physical label at all, and the dynamics
146 model is frozen rather than learned.
- 147 [10] Yuval Tassa, Yotam Doron, Alistair Muldal, Tom Erez, Yazhe Li, Diego de Las Casas, David
148 Budden, Abbas Abdolmaleki, Josh Merel, Andrew Lefrancq, Timothy Lillicrap, and Martin
149 Riedmiller. Deepmind control suite. *arXiv preprint arXiv:1801.00690*, 2018.
- 150 [11] Gaoyue Zhou, Aldo Pacchiano, Ge Yang, Sergey Levine, and Deepak Pathak. Dino-wm:
151 World models on pre-trained visual features enable zero-shot planning. *arXiv preprint*
152 *arXiv:2411.04983*, 2024.